Teams' Self Assesment

***** Heureka Tech Health Framework

Engineering Self-Assessment v1.0

Ø Purpose & Instructions

The **Tech Health Self-Assessment** gives every team a consistent way to evaluate their technical health across **five key areas**.

- Transparent Run once per quarter.
- 1 Led by **Tech Lead / EM** with 1–2 engineers.
- **@** Goal: **reflection & improvement**, not judgment or ranking.
- Scale: 1 = chaotic / ad-hoc, 4 = optimized & automated.
- 🚣 Add short comments ("why this score", "next step").

Scoring Levels (applies to all axes)

Level	Description	Observable Example
1- Chaotic	Ad-hoc, manual, unpredictable.	"Only one person can deploy; tests fail randomly; no docs."
2 - Emerging	Partial coverage or informal discipline.	"Basic tests exist; manual review; occasional postmortems."
3 - Defined	Stable processes & metrics for critical flows.	"Reliable CI; alerting in place; ADRs used for key changes."

4 - Optimized	Automated,	"Pipelines self-validate;
	measurable,	incidents auto-tracked;
	continuously improved.	SLO-based capacity
		steering."

A. TECH DEBT (5 axes)

Sub-Axis	Level 1	Level 2	Level 3	Level 4
Code Quality Debt	No consistent reviews, duplication everywhere. Code inconsistent, high complexity, no shared standards or ownership.	PR reviews exist but no static checks. Manual reviews done inconsistentl y; no linting or static analysis; quality depends on individuals.	Linting & SonarQube in CI; code- owner rules. Consistent PR reviews; automated style and quality checks in CI; ownership defined.	Automated checks + refactoring backlog; clean metrics trend down. Static analysis, code smells tracked; refactoring backlog prioritized; technical debt trend improving (Sonar metrics down).
Architecture / Domain Debt	Monolith / tight coupling. Code and domains are tangled; any change affects	Some domains defined; unclear API contracts. Partial separation, but	Clear bounded contexts; service boundaries reviewed. Domains explicitly	Architecture validated by Fitness Functions; ADR history maintained. Teams own independent

	multiple areas. No clear ownership.	boundaries fuzzy; APIs leak internal logic; unclear data ownership.	defined; APIs stable and versioned; cross- domain dependencie s reviewed regularly.	bounded contexts; minimal coupling; contracts tested automatically and documented.
Infrastructu re Debt	Fully manual process, inconsistent environment s.	Some automation exists, but still requires manual steps or supervision	Reliable CI/CD with rollback.	Full infra-as- code; zero- touch deploy; blue- green or canary releases.
Process / Delivery Debt	Deploys happen anytime, nobody tracks what went out.	Some rules exist (code review, tagging), but people often skip them.	Releases follow a defined flow; everyone knows what's deployed and when.	Deploys fully automated; system tracks DORA metrics automatically .
Knowledge Debt	Context in people's heads.	README outdated; no ADRs.	Docs & ADRs for major services.	Continuous documentati on; onboarding < 1 day.

B. TESTING & AUTOMATION

Sub-Axis	Level 1	Level 2	Level 3	Level 4
Unit Testing	No	Basic unit	Solid coverage of	High-
Coverage	meaningful	tests on key		confidence

	tests; manual validation only	logic; not enforced in CI.	critical code; tests run automatically in CI.	unit suite; fast, stable, and used as living spec.
Integration Testing	No automated integration testing.	Few manual or partial API tests.	Contract tests between main services; checked in CI.	Full API contract validation in pipeline; auto- generated mocks and alerts on contract breaks.
E2E Flow Coverage	Only manual smoke tests after deploy.	Some automated E2E tests on main flows; run nightly.	Key business flows automated and part of CI.	Critical journeys fully automated; regression E2E suite runs on every deploy.
Pipeline Reliability	Builds often fail or give false results.	Pipeline sometimes flaky; unclear root cause.	Reliable CI/CD; success > 95 %, flaky < 5 %.	Self-healing pipelines; > 98 % success; issues visible and auto- reported.
Deployment Automation	Deploys done by hand (SSH, scripts, manual config). No	Basic scripts or CI jobs exist, but require manual input or approvals. Rollback	One-click deploy from CI/CD; rollback tested and documented. No direct	Fully automated delivery pipeline with health checks, canary or

rollback, high	unclear or	server	blue-green
risk of error.	manual.	access	releases, and
		needed.	automatic
			rollback.

○ C. OBSERVABILITY & STABILITY

Sub-Axis	Level 1	Level 2	Level 3	Level 4
Monitoring Coverage	No metrics or alerts; rely on user reports.	Basic infra metrics (CPU/RAM); missing ownership.	Key services monitored with clear owners; SLOs for core SLIs.	Full observability stack (SLI/SLO/Err or Budget) across services; live dashboards used daily by teams.
Alert Hygiene & On-Call Discipline	Alerts noisy or ignored; no on-call setup.	Some alerts tuned; on- call rotation informal or inconsistent.	Alerts actionable (>80%); OpsGenie used for ack/close tracking; P1/P2 severity respected.	On-call fully operational; all alerts actionable; MTTA/MTTR auto-tracked; no ignored pages.
Incident Response (MTTR)	Incidents adhoc; no tracking or ownership.	Incidents logged manually; response unstructured.	MTTR measured; Severity levels defined; RCA (Root Cause	MTTR <1h; incident auto-created in OpsGenie; RCA auto- linked; lessons

			Analyses)req uired for P1s.	shared org- wide.
Logging / Tracing	Scattered logs; no correlation.	Structured logs; tracing partial on key flows.	Centralized logs; tracing active on top paths; searchable incidents.	Full distributed tracing (100%); correlation IDs from edge to DB; logs auto- correlated with incidents.
Postmortem s & Learning	None or blame-based.	Written ad- hoc; no follow-up.	Formal RCAs for major incidents; improvement items tracked.	RCAs for all SEVerity 1– 2s; learnings feed Tech Health backlog; trends reviewed in Tech Health Review.

Metric	Level 1	Level 2	Level 3	Level 4
Deployment Frequency (DF)	< 1 deploy / month.	Weekly deploys.	Daily or per feature branch.	On-demand / continuous delivery.
Lead Time for Changes (LTC)	7 days.	2–7 days.	1–2 days.	< 24 h from commit to prod.

Change Failure Rate (CFR)	25 %.	15–25 %.	5–15 %.	< 5 %.
Mean Time to Recovery (MTTR)	6 h.	2–6 h.	1–2 h.	<1h.
Rollback Frequency	Frequent rollbacks.	Occasional.	Rare (< 10 %).	Exceptional (< 3 %).

© E. GOVERNANCE & KNOWLEDGE

Sub-Axis	Level 1	Level 2	Level 3	Level 4
ADR Discipline	No record of decisions; context lost.	A few ADRs written reactively, not shared.	ADRs created for all significant technical or architectural changes.	ADRs standardized via template; reviewed quarterly; both team- level and global ADRs maintained.
Decision Traceability	No linkage to delivery work (Jira, Epics).	Partial linkage; some ADRs referenced manually.	Decisions traceable to delivery items in ≥80 % of cases where applicable.	Full traceability: each Epic or initiative links to its ADR; org-wide ADRs referenced in standards and playbooks.

Architecture Docs	None or outdated diagrams.	Partial or inconsistent C4 documentati on.	Up-to-date diagrams for core domains; shared in team space.	Comprehensi ve Arc42- style documentati on; automated generation from code/config.
Ownership Clarity	Undefined; "who owns this?" is unclear.	Some ownership lists exist, not maintained.	Clear ownership in Backstage; on-call mapping exists.	Ownership isn't just written down — it's proven by operations. The team that has it on the dashboard and gets the alerts owns it.
Transparenc y	Context closed within teams; decisions siloed.	Shared internally but not visible org-wide.	Visible across engineering; ADRs and docs accessible to all.	"Public by default" culture: decisions and lessons reviewed quarterly in Tech Health Review.

Pulse Survey (Soft Signals – Monthly, 0-10 Scale)

	_
Question	Purpose

1 How much does tech debt slow you down?	Detects friction sources.
2 How smooth is your release process?	Checks delivery flow confidence.
3 Do you have time for stability / engineering work?	Validates 25 % allocation reality.
4 Do you feel leadership supports technical investment?	Measures cultural alignment.
5 Are your tools / environments effective?	Surfaces DevEx blockers.

✓ Interpreting Scores

Average Score	Meaning	Action
1.0 – 1.9	Unstable, reactive.	Create recovery plan; add to Tech Big Rocks.
2.0 - 2.9	Emerging discipline.	Prioritize automation & documentation.
3.0 – 3.4	Stable baseline.	Sustain & optimize critical paths.
3.5 – 4.0	Optimized, data-driven.	Share practices; help mentor others.

Example Scenarios

- **Architecture Debt = Level 2** → legacy monolith; plan for domain decomposition in Q2.
- / Testing = Level 3 → solid CI; next step: flaky-test dashboard.
- Q Observability = Level 1 → no SLOs; urgent stability investment.
- **Delivery = Level 4** → fast CI/CD; monitor CFR trend to keep risk low.
- **@ Governance = Level 2** → ADR discipline inconsistent; add template & quarterly check.



Teams own their improvement roadmap, leadership sees comparable, trendable data, and **tech health becomes a conversation, not a surprise**.